


Parallel implementation of the Gauss-Jordan inversion algorithm using CUDA

*A fast parallel Gauss Jordan algorithm for matrix
inversion using CUDA [1]*

Presentation by:


Sami Hussein Hamid Saeed

ID:31195261



Introduction & Background

- **Why is matrix inversion important?**
- **Why is matrix inversion slow?**
- **Why use GPUs?**
- **Paper Contribution**



Matrix Inversion Uses

Finding the inverse of a matrix can be important for several different reasons such as:

- Structural analyses using finite element method.
- 3D rendering.
- Digital filtering.
- Image processing.
- In general, solving linear equations.

Example:

Given

$$AX = B$$

we can multiply both sides by the inverse of A , provided this exists, to give

$$A^{-1}AX = A^{-1}B$$

But $A^{-1}A = I$, the identity matrix. Furthermore, $IX = X$, because multiplying any matrix by an identity matrix of the appropriate size leaves the matrix unaltered. So

$$X = A^{-1}B$$

Simple example of matrix inversion being used [2].

Gauss-Jordan Algorithm for Matrix Inversion

Until the late 1960's the fastest & simplest method for matrix inversion was the Gauss-Jordan method. It augments a Matrix A and its identity I . Then by performing simple row operations we transform A into I and I becomes A^{-1}

- The method uses a loop to process each column.
- Within each column, the diagonal element ($Matrix[k][k]$) is checked to see whether it is 0 or not, if it is then a swap/addition operation is performed.
- Another loop is run to divide the row k by $matrix[k][k]$ in order to set $matrix[k][k]$ to 1.
- Finally, we must set all elements within column k to 0 by subtracting the rows to negate the values in column k , except the diagonal element at $matrix[k][k]$.
- **Has a worst case time complexity of $O(n^3)$.**
 - Where n is the matrix size.

Algorithm 1 Gauss-Jordan Algorithm

```
1: procedure GAUSSJORDAN( $A$ )
2:   Let  $A$  be the input matrix
3:   Let  $I$  be the identity matrix of the same size as  $A$ 
4:   for  $k = 1$  to  $n$  do
5:     Find the row  $r$  such that  $A[r][k]$  is the maximum among  $A[k][k]$ ,
6:      $A[k+1][k]$ , ...,  $A[n][k]$ 
7:     Swap rows  $A[r]$  and  $A[k]$ 
8:     Divide row  $A[k]$  by  $A[k][k]$ 
9:     Divide row  $I[k]$  by  $A[k][k]$ 
10:    for  $i = 1$  to  $n$  do
11:      if  $i$  is not equal to  $k$  then
12:         $factor \leftarrow A[i][k]$ 
13:        for  $j = 1$  to  $n$  do
14:           $A[i][j] \leftarrow A[i][j] - factor \times A[k][j]$ 
15:           $I[i][j] \leftarrow I[i][j] - factor \times I[k][j]$ 
16:        end for
17:      end if
18:    end for
19:     $A^{-1}$  is the matrix  $I$ 
20: end procedure
```

Why GPUs?

- The Gauss-Jordan method is suitable for massive parallelisation.
 - Lots of independent operations performed sequentially. Each row has N elements where row operations can be done on all elements in parallel.
- GPU architecture supports massively parallel processing by enabling the usage of thousands of threads in parallel.
 - GPUs generally have much more ALUs than a CPU.
- CPUs only allow for 8 to 12 threads.
- Thread creation and memory transfer cost is much lower in GPU.




Comparison of CPU and GPU architecture [3]

The paper's proposed contributions are:

- **Compare** and **contrast** different matrix inversion methods and algorithms.
- **Redesign** the Gauss-Jordan algorithm in order to exploit massively multithreaded GPUs.
- Perform **testing** on different **types** of matrices of different **sizes**.
- Prove that the time complexity of the algorithm can scale as **$O(n)$** , **provided n^2 threads** can be created.

Hypothesis & Problem Statement

- **Hypothesis**
- **Problem Statement**



Problem Statement

Most of the **basic** matrix inversion methods are usually done by simply performing **repeated independent row operations**. Each row operation requires a loop to iterate through each column in a row.


Problems faced in matrix inversion:

- All of the row operations need to be performed on each element within a row, meaning that at best, a serial implementation of $O(n^2)$ might be possible in the future.
- A large number of independent row and column operations are done in a sequential order.
- CPU parallelism is insufficient for larger matrices as it can only support a limited number of threads.

Hypothesis

The hypothesis of the paper is that the proposed parallel Gauss-Jordan inversion algorithm implemented on CUDA can significantly reduce the worst case time complexity for matrix inversion from $O(n^3)$ to a linear $O(n)$ provided n^2 threads.

Related work



Related work

Methods	Pros	Cons
Strassen	<ul style="list-style-type: none"> - First algorithm to reduce matrix inversion from $O(n^3)$ to $O(n^{2.808})$. - This was done by recursively breaking down the matrices into smaller submatrices and performing fewer arithmetic operations. - Has been repeated improved by different researchers using different methods to test the lowest limit of exponent value, which is suggested to be 2. 	<ul style="list-style-type: none"> - Cannot be easily parallelised due to its recursive nature. - Requires a large value of N to be parallelised. - Accuracy of the inverse depends on the sub-matrix selected in the initial step.
Strassen-Newton	<ul style="list-style-type: none"> - A parallel variation of strassen's method using newton iterations. - Can yield moderate results even on a single CPU. - Up to 55% speedup can be observed even on a single processor. 	<ul style="list-style-type: none"> - Design is coupled with the limitations of the time's architecture (Published in 1988). - Speedup is relatively low. - Does not take into account the usage of GPUs.
Coppersmith & Winograd	<ul style="list-style-type: none"> - Further improves matrix inversion using strassen's new laser method. - Avoids unnecessary arithmetic operations. - Considerably improved the runtime complexity to a record breaking $O(n^{2.376})$. 	<ul style="list-style-type: none"> - Only applicable to specific types of matrices - Extremely complex to implement without a deep understanding of linear algebra.
LUP Decomposition	<ul style="list-style-type: none"> - Decomposes the matrix into three sub matrices, a lower triangular matrix, upper triangular matrix and a permutation matrix. - More efficient when dealing with a large value of N. - Numerically stable. 	<ul style="list-style-type: none"> - Consumes a lot of RAM to store all 3 matrices. - Time complexity is still non-linear.


Related work

Methods	Pros	Cons
Cholesky Decomposition	<ul style="list-style-type: none"> - Also performs decomposition of the matrix into submatrices, a lower triangular matrix and its transpose. - Further optimizes the algorithm for specific matrices such as symmetric matrices and matrices with positive definite. 	<ul style="list-style-type: none"> - Consumes a lot of RAM to store 2 matrices. - Only applicable to select matrices. - Tough to parallelise due to dependencies.
QR Decomposition	<ul style="list-style-type: none"> - Breaks down the matrix into an orthogonal matrix Q and a upper triangular matrix R. - Numerically stable. - Can be applied to rectangular matrices. 	<ul style="list-style-type: none"> - Consumes a lot of RAM to store 2 matrices. - Requires additional pre-processing for some irregular matrices.
RRQR Factorization	<ul style="list-style-type: none"> - Further improvement of the QR Decomposition method. - Focuses on more important parts of the matrix. Thus allowing for less computations do be done. 	<ul style="list-style-type: none"> - Inherits the same space and pre-processing limitations from QR decomposition.
Monte Carlo Methods	<ul style="list-style-type: none"> - Provides an approximate solution using random sampling. - No set time complexity. - Used for inverting the Hermitian matrix and positive definite matrix. 	<ul style="list-style-type: none"> - Subject to sampling errors. - Converging to true solution may require large amount of samples. - Only applicable to specific matrices.

Common limitations found in most papers:

- Most papers aim to reduce n 's exponent by a small fraction.
- Most papers do not address the use of massive parallelisation.
 - This is mainly because all of the **advanced** optimizations made in these papers require some form of **dependency** between each iteration, making the code **non parallelizable**.
- Some of the fastest methods are only applicable for specific type of matrices.

Methodology



Methodology - Approach

- 1) **Outermost loop to iterate through each column.**
 - a) Cannot be done in parallel due to dependency between iterations.
- 2) **Ensure that the diagonal element ($\text{matrix}[j][j]$) is non-zero.**
 - a) Find row k where column j is not 0.
 - b) Add the two rows together in parallel (1 thread for each column).
- 3) **Convert the diagonal element into 1 by dividing the entire row by $\text{matrix}[j][j]$**
 - a) The division of the entire row can be done in parallel using 1 thread to perform the division on each column.
- 4) **Convert all other rows in the column to 0**
 - a) This is done by simply subtracting the value by itself multiplied by the row containing the diagonal.
 - b) There are n blocks created containing n threads each. A **block** represents a **column** and each **thread** represents a **row** for each column/block.

```
Read matrix
Initialize  $n$  to size of matrix
Initialize  $j$  to 0
while  $j < n$ , do:

    Find  $k$  where  $\text{matrix}[k][j]$  is not 0
    Spawn  $n$  threads in 1 block
    for thread  $i$  of  $n$  in block 1, do:
         $\text{matrix}[j][i] = \text{matrix}[j][i] + \text{matrix}[k][i]$ 
    end for

    Spawn  $n$  threads in 1 block
    for thread  $i$  of  $n$  in block 1, do:
         $\text{matrix}[j][i] = \text{matrix}[j][i] / \text{matrix}[j][j]$ 
    end for

    Spawn  $n$  threads each in  $n$  blocks
    for thread  $i$  of  $n$  in block  $r$ , do:
         $\text{matrix}[i][r] = \text{matrix}[i][r] - \text{matrix}[i][j] * \text{matrix}[j][r]$ 
    end for
    Increase  $j$  by 1

end while
Write matrix
```


Methodology - CUDA Implementation

Row Fixing

- Each row must be divided by its diagonal element in order to turn the diagonal element itself into 1.
- Here this is done by multithreading, where each thread represents a column and all divisions are done in parallel

$$\begin{bmatrix}
 1 & \cdots & a_{1j} & a_{1(j+1)} & \cdots & a_{1n} & a_{11}^{inv} & \cdots & 0 & 0 & \cdots & 0 \\
 0 & \cdots & a_{2j} & a_{2(j+1)} & \cdots & a_{2n} & a_{21}^{inv} & \cdots & 0 & 0 & \cdots & 0 \\
 \vdots & \cdots & \vdots & \vdots & \cdots & \vdots & \vdots & \cdots & \vdots & \vdots & \cdots & \vdots \\
 0 & \cdots & a_{jj} & a_{j(j+1)} & \cdots & a_{jn} & a_{j1}^{inv} & \cdots & 1 & 0 & \cdots & 0 \\
 \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\
 0 & \cdots & a_{nj} & a_{n(j+1)} & \cdots & a_{nn} & a_{n1}^{inv} & \cdots & 0 & 0 & \cdots & 1
 \end{bmatrix}$$

n elements processed

```

__global__ void fixRow(float *matrix, int size, int rowId) {
    // The ith row of the matrix
    __shared__ float Ri[512];
    // The diagonal element for ith row
    __shared__ float Aii;
    int colId = threadIdx.x;
    Ri[colId] = matrix[size * rowId + colId];
    Aii = matrix[size * rowId + sharedRowId];
    __syncthreads();
    // Divide the whole row by the diagonal element making sure it is not 0
    Ri[colId] = Ri[colId] / Aii;
    matrix[size * rowId + colId] = Ri[colId];
}
    
```

Column Fixing

- Each row in each column must be subtracted by itself multiplied by the diagonal element in order to turn the entire column of the diagonal element into 0.
- Here this is done by multithreading, where each block represents a column and each thread within the block represents a row and then all subtractions are done in parallel.

$$\begin{bmatrix}
 1 & \cdots & a_{1j} & a_{1(j+1)} & \cdots & a_{1n} & a_{11}^{inv} & \cdots & 0 & 0 & \cdots & 0 \\
 0 & \cdots & a_{2j} & a_{2(j+1)} & \cdots & a_{2n} & a_{21}^{inv} & \cdots & 0 & 0 & \cdots & 0 \\
 \vdots & \cdots & \vdots & \vdots & \cdots & \vdots & \vdots & \cdots & \vdots & \vdots & \cdots & \vdots \\
 0 & \cdots & 1 & a_{j(j+1)} & \cdots & a_{jn} & a_{j1}^{inv} / a_{jj} & \cdots & 1/a_{jj} & 0 & \cdots & 0 \\
 \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\
 0 & \cdots & a_{nj} & a_{n(j+1)} & \cdots & a_{nn} & a_{n1}^{inv} & \cdots & 0 & 0 & \cdots & 1
 \end{bmatrix}$$

n^2 elements
 processed


```

__global__ void fixColumn(float *matrix, int size, int colId) {
    int i = threadIdx.x;
    int j = blockIdx.x;
    // The colId column
    __shared__ float col[512];
    // The jth element of the colId row
    __shared__ float AColIdj;
    // The jth column
    __shared__ float colj[512];
    col[i] = matrix[i * size + colId];
    if (col[i] != 0) {
        colj[i] = matrix[i * size + j];
        AColIdj = matrix[colId * size + j];
        if (i != colId) {
            colj[i] = colj[i] - AColIdj * col[i];
        }
        matrix[i * size + j] = colj[i];
    }
}
    
```

Overall the methodology addresses the two main issues identified earlier since:

- 1) The proposed approach indicates, that a linear runtime complexity algorithm for all matrix types is possible since the new method **has parallelised all the loops inside the main loop.**
- 2) By using an architecture that supports **massively parallel** execution such as GPUs we can effortlessly and quickly create a large amount of threads compared to the limited and slow CPU threads.

Analysis of Results & Conclusions



The testing system uses a Intel Quad Core processor Q8400 @ 2.66 GHz each. The GPU used for executing the CUDA code is a GTX 260.

The authors conduct the runtime tests using different types of matrices:

- Sparse Matrix
 - A matrix where most elements are 0.
- Band Matrix
 - A sparse matrix where it's non-zero elements are located in the diagonal region.
- Hollow Matrix
 - A matrix where the diagonal elements of a matrix are 0. This is the **most computationally expensive type of matrix to inverse**, as it requires row addition or swapping to be done for each column.
- Identity Matrix
 - A matrix where all the elements are set to 0, except the diagonal elements, which are set to 1
- RandomMatrix

Results

- All GPU tests show that for almost all matrices the runtime scales linearly with matrix size.
- There first quadratic curvature can be noticed at around $n=100$, due to the number of threads requested being more than the thread available in a GTX 260.
- Hollow matrix has a steeper gradient due to it being the matrix with the most number of computations necessary.
- Overall, the code performs almost 10 times as fast in a GPU than a CPU.

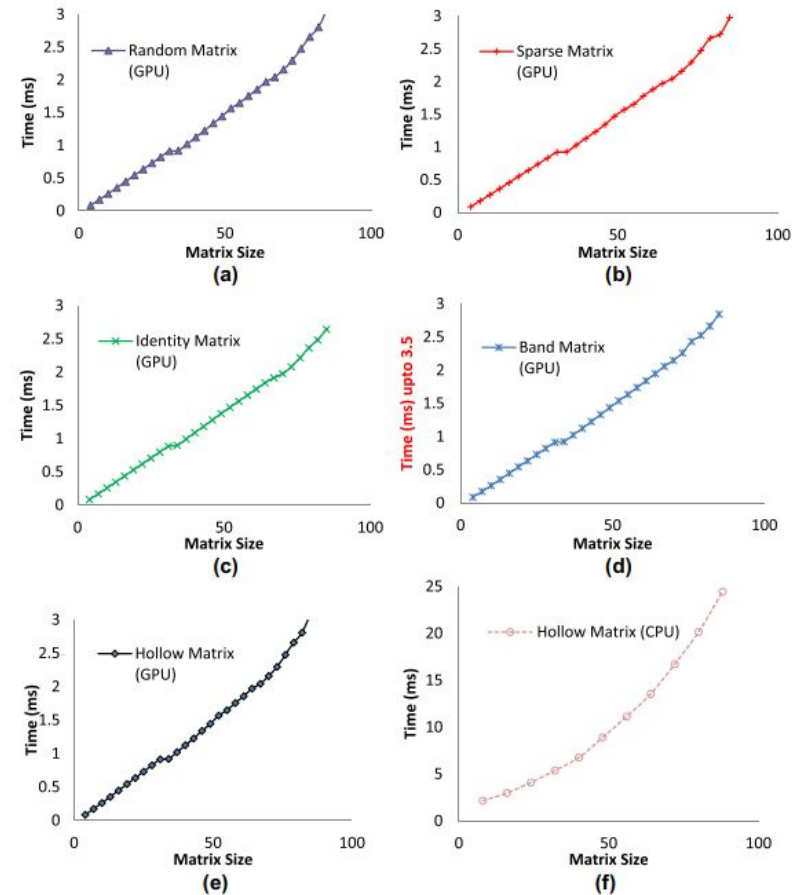


Fig. 4. (a-e): Linear computation time for matrix inversion is observed up to $n \approx 64$ using GPU, (f) computation time for inverting hollow matrix using CPU.

Results (Scaled Up)

- GPU code no longer exhibits a linear increase in runtime as the matrix size increases.
- The change in the time complexity can be attributed to the fact that the GTX 260 GPU can store 9216 threads and 288 dispatched threads concurrently.
- Given the need for 262,144 threads to process a 512-sized matrix, the GPU's thread capacity is exceeded, necessitating some threads to be executed sequentially.
- GPU code is still considerably faster than the CPU code, regardless of the size.

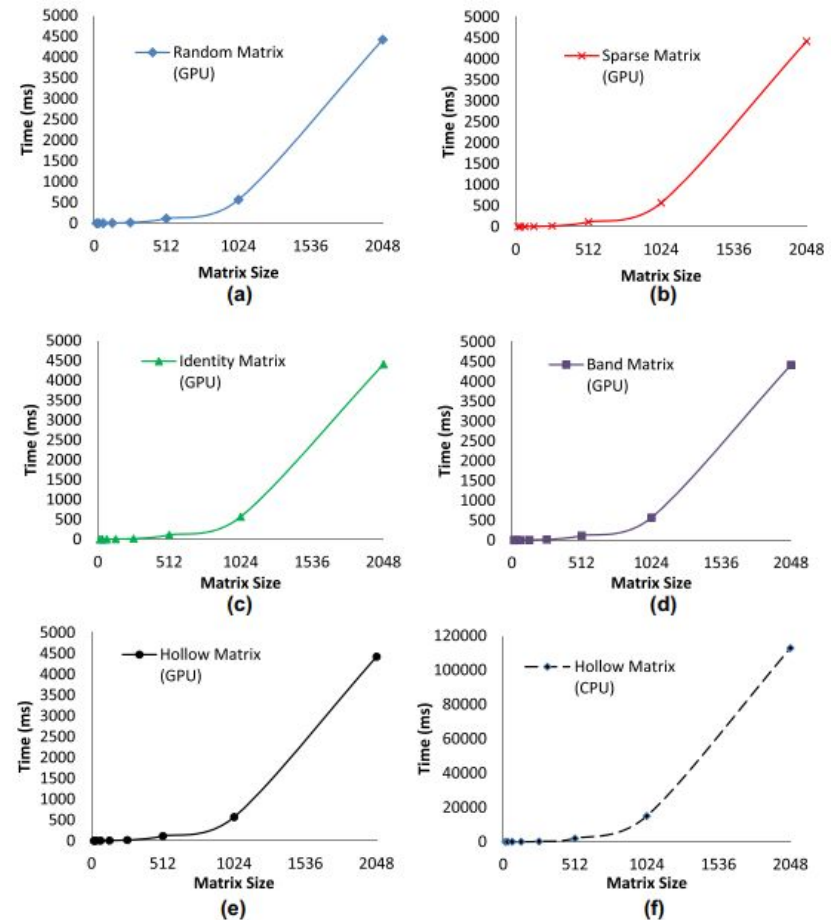


Fig. 5. Computation time for inverting different types of matrices, (a-e) using GPU, (f) using CPU.

Overall, the tests performed by the researchers show that the tests do indeed support the hypothesis made, by showing that:

- 1) The worst case time complexity of the matrix inversion does in fact stay at $O(n)$ when allocated n^2 threads.
- 2) The algorithm can work for any type of matrix.

References

- [1] G. Sharma, A. Agarwala, and B. Bhattacharya, "A fast parallel Gauss Jordan algorithm for matrix inversion using CUDA," *Computers & Structures*, vol. 128, pp. 31–37, Nov. 2013, doi: 10.1016/j.compstruc.2013.06.015.
- [2] Math Centre. "Engineering Math First Aid Kit: Using inverse matrix to solve equations" Math Centre, n.d., https://www.mathcentre.ac.uk/resources/Engineering%20maths%20first%20aid%20kit/latexsource%20and%20diagrams/5_6.pdf.
- [3] Thambawita, V., Vajira, R., Ragel, R., & Elkaduwe, D. (2014). "To Use or Not to Use: Graphics Processing Units for Pattern Matching Algorithms." IEEE.